

Assignment / Explore Query Planning and Indexing

Summer Full 2023

Setup

Creating a connection variable to the SQLite database to perform all the query operation

```
library(RSQLite)

#provide path to database folder

#provied path to database file

dbfile = "sakila.db"

dbcon <- dbConnect(RSQLite :: SQLite() , dbfile)

tables <- dbListTables(dbcon)

# Print the table names
print(tables)
```



```
## [1] "actor"           "actor_info"
## [3] "address"         "category"
## [5] "city"            "country"
## [7] "customer"        "customer_list"
## [9] "film"             "film_actor"
## [11] "film_category"   "film_list"
## [13] "film_text"       "inventory"
## [15] "language"        "nicer_but_slower_film_list"
## [17] "payment"         "rental"
## [19] "sales_by_film_category" "sales_by_store"
## [21] "sqlite_sequence" "staff"
## [23] "staff_list"      "store"
```

Ques 1

Below is the R chunk to perform deletion of all user defined indexes and query to find the number of films per language

```
# TODO : Delete user defined indexes

drop_ind <- "DROP INDEX IF EXISTS TitleIndex"

dbExecute(dbcon , drop_ind)
```

```
## [1] 0
```

```
query <- "SELECT l.name AS language_name , COUNT(*) AS Total_films
FROM film as f
JOIN language as l ON f.language_id = l.language_id
GROUP BY l.name"
```

```
df <- dbGetQuery(dbcon, query)

print(df)
```

```
##   language_name Total_films
## 1      English      1000
```

Ques 2

Below is the R chunk to find the query plan for above query

```
query2 <- " EXPLAIN QUERY PLAN
SELECT l.name AS language_name , COUNT(*) AS Total_films
FROM film as f
JOIN language as l ON f.language_id = l.language_id
GROUP BY l.name"
```

```
result2 <- dbGetQuery(dbcon, query2)

print(result2)
```

```
##   id parent notused          detail
## 1  7      0      0          SCAN f
## 2  9      0      0 SEARCH 1 USING INTEGER PRIMARY KEY (rowid=?)
## 3 12      0      0          USE TEMP B-TREE FOR GROUP BY
```

Ques 3

Below is the R chunk to find the title , category name and length of the film, titled “ZORRO ARK”

```
query3 <- "SELECT f.title , f.length , c.name
FROM film as f
JOIN film_category as fc ON f.film_id = fc.film_id
JOIN category as c ON fc.category_id = c.category_id
where f.title = 'ZORRO ARK'
"

# calculating time taken here for the purpose of ques 8
s_t <- system.time ({
  result3 <- dbGetQuery(dbcon, query3)
})
```

```
print(s_t)
```

```
##      user  system elapsed  
##    0.001    0.001    0.001
```

```
print(result3)
```

```
##      title length  name  
## 1 ZORRO ARK     50 Comedy
```

Ques 4

Below is the R chunk to explain the query plan of above query

```
query4 <- " EXPLAIN QUERY PLAN  
SELECT f.title , f.length , c.name  
FROM film as f  
JOIN film_category as fc ON f.film_id = fc.film_id  
JOIN category as c ON fc.category_id = c.category_id  
where f.title = 'ZORRO ARK'  
"
```

```
result4 <- dbGetQuery(dbcon, query4)
```

```
print(result4)
```

```
##      id parent notused  
## 1  4      0      0  
## 2  6      0      0  
## 3  9      0      0  
##                                     detail  
## 1 SCAN fc USING COVERING INDEX sqlite_autoindex_film_category_1  
## 2      SEARCH c USING INTEGER PRIMARY KEY (rowid=?)  
## 3      SEARCH f USING INTEGER PRIMARY KEY (rowid=?)
```

Ques 5

Below is the R chunk to create index film table

```
dbExecute(dbcon, "CREATE INDEX TitleIndex ON FILM (TITLE)")
```

```
## [1] 0
```

Ques 6

Below is the R chunk to create index film table

```

query6_plan <- " EXPLAIN QUERY PLAN
SELECT f.title , f.length , c.name
FROM film as f
JOIN film_category as fc ON f.film_id = fc.film_id
JOIN category as c ON fc.category_id = c.category_id
where f.title = 'ZORRO ARK'
"

result6_plan <- dbGetQuery(dbcon, query6_plan)

print(result6_plan)

```

```

##      id parent notused
## 1   5      0      0
## 2  10      0      0
## 3  14      0      0
##
##                                     detail
## 1                                     SEARCH f USING INDEX TitleIndex (title=?)
## 2 SEARCH fc USING COVERING INDEX sqlite_autoindex_film_category_1 (film_id=?)
## 3                                     SEARCH c USING INTEGER PRIMARY KEY (rowid=?)

```

Ques 7

Below is the R chunk to explain the index working :

Difference : (6) searches the films table by the index of TitleIndex where as (4) uses primary key (rowid) to search

How to identify : The search from index is specified by index name, here in case of (6) we have “SEARCH f USING INDEX TitleIndex (title=?)”

Comments : where there is a search by title query having an index in title column of the table ensures the db engine improve the efficiency of data retrieval operations by a sorted copy of the indexed column(s) and corresponding data rows.

QUES 8

We can see that after creating index the operation happened relatively fast, as it incorporated efficient search with sorted index rather than linear scan

```

time <- system.time({
  film_info <- dbGetQuery(dbcon, "
    SELECT f.title, c.name AS category, f.length
    FROM film f
    JOIN film_category fc ON f.film_id = fc.film_id
    JOIN category c ON fc.category_id = c.category_id
    WHERE f.title = 'ZORRO ARK'
  ")
})

```

```
print(time)
```

```
##      user  system elapsed
##        0        0        0
```

QUES 9

*# as per readings using lower case because MySQL queries are not case-sensitive
by default. If this works different it might be because of the internals of
engine*

```
query9 <- "  
  SELECT f.title , f.length , l.name  
  FROM film as f  
  JOIN language as l ON f.language_id = l.language_id  
  where f.title LIKE '%gold%'  
  "
```

```
result9 <- dbGetQuery(dbcon, query9)
```

```
print(result9)
```

```
##              title length   name  
## 1      ACE GOLDFINGER     48 English  
## 2  BREAKFAST GOLDFINGER    123 English  
## 3           GOLD RIVER    154 English  
## 4 GOLDFINGER SENSIBILITY    93 English  
## 5      GOLDMINE TYCOON    153 English  
## 6           OSCAR GOLD    115 English  
## 7  SILVERADO GOLDFINGER     74 English  
## 8           SWARM GOLD    123 English
```

*# as per readings using lower case because MySQL queries are not case-sensitive
by default. If this works different it might be because of the internals of
engine*

```
query10 <- "  
  EXPLAIN QUERY PLAN  
  SELECT f.title , f.length , l.name  
  FROM film as f  
  JOIN language as l ON f.language_id = l.language_id  
  where f.title LIKE '%gold%'  
  "
```

```
result10 <- dbGetQuery(dbcon, query10)
```

```
print(result10)
```

```
##   id parent notused      detail  
## 1  3      0        0      SCAN f  
## 2  8      0        0 SEARCH 1 USING INTEGER PRIMARY KEY (rowid=?)
```

The query doesn't use indexes because : Indexes are not useful when pattern matching searches are performed with LIKE. And in the case of searches using LIKE the engines do linear scan anyways.